

# Introduction to Operating Systems

Troels Henriksen

Incorporating material by Mathias Payer  
(<https://github.com/HexHive/OSTEP-slides>).

# The purpose of operating systems

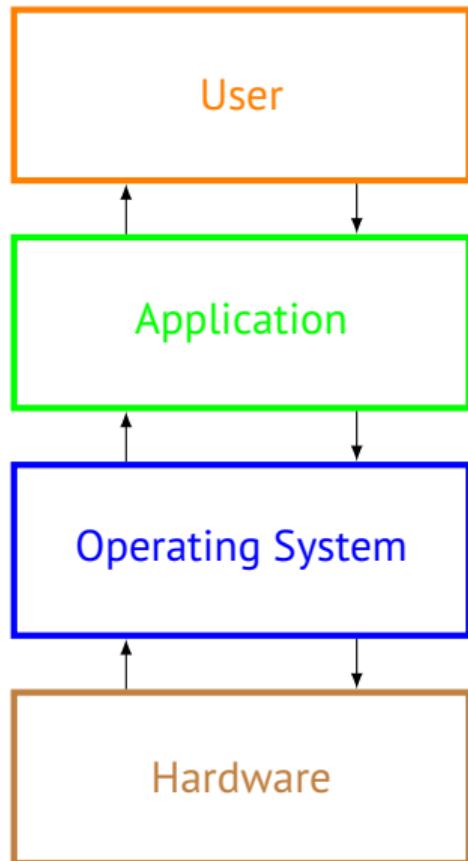
Kernel and Processes

System calls

Process management in Unix

Main takeaways

# What is an operating system?



OS is middleware between applications and hardware.

- Provides standardized interface to resources.
- Manages hardware.
- Orchestrates currently executing processes.
- Responds to resource access requests.
- Handles access control.

# Why study operating systems?

## Inspirational

- One of the most potent *abstractions* in computing.
  - ▶ Each process thinks it has machine to itself.
  - ▶ Controlled communication.
  - ▶ Abstracts over hardware differences.

## Practical

- You almost always use an operating system.
- Its performance characteristics are important to understand.
- It often determines what is fundamentally possible.

**They are where the magic happens.**

# In the old days

Each brand of machine would have its own operating system.



IBM System/360 running OS/360 (man not part of computer)



DEC PDP-10 running TENEX



VAX 11/780 running VMS



- PDP-11 running early UNIX, written in C.
- Developed at AT&T, who were banned from selling UNIX.
- Shared it (almost) freely with others, who *ported* UNIX to every machine under the sun.
- Unix was *popular, good enough, and cheap.*

# What is Unix

- What?**
  - Unix is a family of operating systems derived from original UNIX developed in the 1970s by Ken Thompson and Dennis Ritchie.
  - Most modern operating systems are heavily influenced by Unix (even Windows).
  - Many operating systems are *direct descendants*: Linux, iOS, macOS, the BSDs, etc.
- Why?**
  - We teach Unix because it is *simple* and *representative* of modern systems.
  - We use Unix designs for all examples.
  - ...does not mean Unix is always *good* design.

The purpose of operating systems

Kernel and Processes

System calls

Process management in Unix

Main takeaways

# Processes contra programs

**Program:** is a file containing code. Stateless and **dead**.

**Process:** a running *instance* of a program. The same program can be running in multiple instances. Stateful and **alive**.

**They are not the same thing, although we informally often say *program* when we really mean *process*.**

# Processes contra programs

**Program:** is a file containing code. Stateless and **dead**.

**Process:** a running *instance* of a program. The same program can be running in multiple instances. Stateful and **alive**.

**They are not the same thing, although we informally often say *program* when we really mean *process*.**

- Operating systems manage *processes*.
  - ▶ Switching between multiple *concurrent* processes.
  - ▶ Handling process termination.
  - ▶ Starting new processes (perhaps from a given *program file*).

# The kernel

- Technically, *operating systems* encompass lots of parts: shell, GUI, C library, maybe bundled applications (mail reader etc).
- In this course, when we say *operating system* we really mean the **kernel**.

# The kernel

- Technically, *operating systems* encompass lots of parts: shell, GUI, C library, maybe bundled applications (mail reader etc).
- In this course, when we say *operating system* we really mean the **kernel**.

## Kernel

Always-resident code that services requests from the hardware and manages processes. **It is not a process.**

- The kernel uses the same CPU, memory, and other hardware as ordinary code.
- When running process code, the CPU is in *unprivileged state*, and many operations are restricted (e.g. access to hardware devices).
  - ▶ When an *interrupt* happens, the CPU switches to *privileged state* and jumps to kernel code, which handles it and then resumes the previously running process.
    - ▶ Think of it like a sudden and unplanned procedure call.
    - ▶ Interrupts can be outside events (keyboard press, network traffic) or special instructions (invalid memory accesses, *system calls*).

# Virtualising the CPU

**Goal** Give each process the illusion of exclusive CPU access.

**Reality:** the CPU is a shared resource.

# Virtualising the CPU

**Goal** Give each process the illusion of exclusive CPU access.

**Reality:** the CPU is a shared resource.

## Solution: **context switching**

- Only one process gets to run at a time.
- ...but we regularly switch between available processes.
- Doing this often and rapidly creates the illusion of simultaneous execution.

# Context switching

**Intuition** Pausing a process, saving its entire *state*, then resuming some other process based on its saved state.

# Context switching

**Intuition** Pausing a process, saving its entire *state*, then resuming some other process based on its saved state.

## So what do we need to save?

1. All registers, including control registers.
2. Contents of memory.

# Context switching

**Intuition** Pausing a process, saving its entire *state*, then resuming some other process based on its saved state.

## So what do we need to save?

1. All registers, including control registers.
2. Contents of memory.

## So when do we do this?

- Regular *timer interrupts* transfer control to the kernel, whose *scheduler* decides the next process to run.
  - ▶ Scheduling is a big and interesting topic that we don't have time to go into.

The purpose of operating systems

Kernel and Processes

System calls

Process management in Unix

Main takeaways

# System calls

- Only the kernel has direct access to hardware and system memory.
- Whenever we want to do IO we have to perform a *system call*.

## System calls

A request by a process that the kernel carries out some operation on its behalf.

- Much like a function call, but implemented very differently.

# System calls in RISC-V

- The `ecall` (*environment call*) instruction transfers control to the kernel.
  - ▶ Kernel then inspects registers (mostly `a0-a7`) to see what it has been asked to do.
  - ▶ Specific interpretation varies between operating systems.
  - ▶ System call identified by a number.

# System calls in RISC-V

- The `ecall` (*environment call*) instruction transfers control to the kernel.
  - ▶ Kernel then inspects registers (mostly `a0-a7`) to see what it has been asked to do.
  - ▶ Specific interpretation varies between operating systems.
  - ▶ System call identified by a number.

## In RARS

<https://github.com/TheThirdOne/rars/wiki/Environment-Calls>

- System call number passed in `a7`.
- Excerpt:

Human name	Number	Description	Reads	Writes
PrintInt	1	Prints int to console.	a0	
ReadInt	5	Reads int from console.		a0

## Used in A0's io.s

```
read_loop:  
    beq    t1 , t2 , read_done  
    li     a7 , 5  
    ecall           # read int  
    sw    a0 , 0(t0)  
    addi   t0 , t0 , 4      # next output addr  
    addi   t1 , t1 , 1      # increment count  
    jal    zero , read_loop
```

## System calls in C

- C exposes system calls as functions.
  - ▶ Internally use `ecall` instruction or architecture-specific equivalent.
- Not observably distinct from other C functions, except typically more primitive and tedious to use.

# System calls in C

- C exposes system calls as functions.
  - ▶ Internally use `ecall` instruction or architecture-specific equivalent.
- Not observably distinct from other C functions, except typically more primitive and tedious to use.

## Example

```
// system calls
int open(const char *pathname, int flags);
ssize_t write(int fd, const void *buf, size_t count);

// stdio functions
FILE *fopen(const char *pathname, const char *mode);
size_t fwrite(const void *ptr, size_t size, size_t nmemb,
              FILE *stream);
```

# File descriptors

```
int open(const char *pathname, int flags);  
  
ssize_t write(int fd, const void *buf, size_t count);
```

- The `int` returned by `open` is a *file descriptor*.
- Has no significance in itself, but allows the kernel to recognise the open file when passed to other system calls.
  - ▶ Typically an index into some kernel-side table.
  - ▶ Such values are known as *handles*.
- Passing complex data structures or pointers between *kernel space* and *user space* is annoying and fragile, so we usually use numeric identifiers instead.

The purpose of operating systems

Kernel and Processes

System calls

Process management in Unix

Main takeaways

# Basic principles

- Each process in Unix has a *process ID* (PID).
- Each process has a *parent*.
  - ▶ ...except the *initial process* (`init`) with PID 1.
- A process may have multiple *children*.
- Implies processes are organised as a *tree* (`pstree` command shows it).
- **Creating processes:** `fork()`.
- **Terminating current process:** `exit()`.
- **Loading program code from disk into current process:** `exec()`.
- **Waiting for a specific child to die:** `waitpid()`.
- **Getting PID of running process:** `getpid()`.

The purpose of operating systems

Kernel and Processes

System calls

Process management in Unix

Main takeaways

## Main takeaways

- Hardware provides *mechanisms* such as interrupts, privileged/unprivileged mode, and *virtual memory* (next week).
  - ▶ Kernels implement *policy* and *abstractions* on top.
- Processes are a *purely virtual concept*—CPU has no idea what they are.
- Processes are *isolated* from each other.
- Processes can only directly interact with the outside world through *system calls*, mediated by the kernel.